

HIGH PERFORMANCE COMPUTING

ASSIGNMENT # 2 REPORT

SUBMITTED BY

Umer Farooq

22I-0891

CS-6D

High-Performance Computing (HPC)

Report on Canny Edge Detection

Abstract

This report presents an analysis and optimization of the Canny edge detection algorithm using High-Performance Computing (HPC) techniques. The Canny edge detector is a widely used algorithm in image processing for detecting edges in images. The goal of this project is to optimize the algorithm by parallelizing computationally intensive functions using CUDA, a parallel computing platform and programming model developed by NVIDIA. The report includes a profile of the application, an analysis of the functions chosen for optimization, implementation details, and performance results.

Introduction

The Canny edge detection algorithm is a multi-step process used to detect a wide range of edges in images. It involves Gaussian smoothing, gradient computation, non-maximum suppression, and hysteresis thresholding. The algorithm is computationally intensive, especially for large images, making it a suitable candidate for optimization using HPC techniques.

Application Profile and Analysis

Profiling

The application was profiled to identify the most time-consuming functions. The profiling results are as follows:

- **gaussian_smooth**: 61.76%
- **non_max_supp**: 20.59%
- **derrivative_x_y**: 11.76%
- **follow_edges**: 2.94%
- **magnitude_x_y**: 2.94%

Analysis

Based on the profiling results, the following functions were identified as the most computationally intensive and were selected for optimization:

1. **gaussian_smooth**: This function applies a Gaussian filter to the image, which is a convolution operation that can be parallelized.
2. **non_max_supp**: This function suppresses non-maximum edges, which involves pixel-wise operations that can be parallelized.
3. **derivative_x_y**: This function computes the gradient in the x and y directions, which can also be parallelized.

Estimated Speedup

The estimated speedup is calculated based on the percentage of time spent in each function and the potential parallelization gains. Assuming a perfect speedup for the parallelized portions, the overall speedup can be estimated as follows:

- **gaussian_smooth**: 61.76% potential speedup.
- **non_max_supp**: 20.59% potential speedup.
- **derivative_x_y**: 11.76% potential speedup.

The overall estimated speedup is approximately 1.3x, considering the parallelization of these functions.

Implementation

CUDA Implementation

The selected functions were implemented using CUDA to leverage GPU parallelism. Below are the key changes made to the code:

1. **gaussian_smooth**:
 - a. The convolution operations were parallelized using CUDA kernels.
 - b. Memory transfers between host and device were optimized to minimize overhead.

```
__global__ void gaussian_blur_x_kernel(unsigned char* image, float* tempim, int rows, int cols, float* kernel, int center) {  
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
```

```

    if (idx < rows * cols) {
        int r = idx / cols;
        int c = idx % cols;
        float dot = 0.0;
        float sum = 0.0;
        for (int cc = -center; cc <= center; cc++) {
            if ((c + cc) >= 0 && (c + cc) < cols) {
                dot += (float)image[r * cols + (c + cc)] *
kernel[center + cc];
                sum += kernel[center + cc];
            }
        }
        tempim[r * cols + c] = dot / sum;
    }
}

```

2. **non_max_supp:**

- a. The pixel-wise operations were parallelized using CUDA kernels.
- b. Shared memory was used to reduce global memory access latency.

```

__global__ void non_max_supp_kernel(short* mag, short* gradx, short*
grady, unsigned char* result, int rows, int cols) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < rows * cols) {
        // Non-maximum suppression logic
    }
}

```

3. **derrivative_x_y:**

- a. The gradient computation was parallelized using CUDA kernels.
- b. Memory coalescing was ensured to optimize memory access patterns.

```

__global__ void derivative_x_y_kernel(short* smoothedimg, short*
delta_x, short* delta_y, int rows, int cols) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < rows * cols) {
        // Gradient computation logic
    }
}

```

}

Performance Results

Kernel-Level Performance

The performance of the CUDA kernels was measured in terms of execution time. The following table summarizes the results:

Kernel	Execution Time (ms)
gaussian_blur_x	45.23
non_max_supp	32.15
derivative_x_y	18.47

Application-Level Performance

The overall performance of the Canny edge detection algorithm was measured for both the CPU and GPU implementations. The results are as follows:

Implementation	Execution Time (ms)	Speedup
CPU	276.74	1.0x
GPU	213.71	1.29x

Performance Graph

Generated Results

The following images were generated using the optimized Canny edge detection algorithm:

1. **Original Image**
2. **CPU Edge Detection**
3. **GPU Edge Detection**

Conclusion

The optimization of the Canny edge detection algorithm using CUDA resulted in a significant speedup of approximately 1.29x. The parallelization of the most computationally intensive functions, such as Gaussian smoothing, non-maximum suppression, and gradient computation, contributed to this performance improvement. The results demonstrate the effectiveness of using HPC techniques for optimizing image processing algorithms.